

Introduction to Linux Kernel Programming

CS423 2016 Spring
Haitong Tian



ILLINOIS

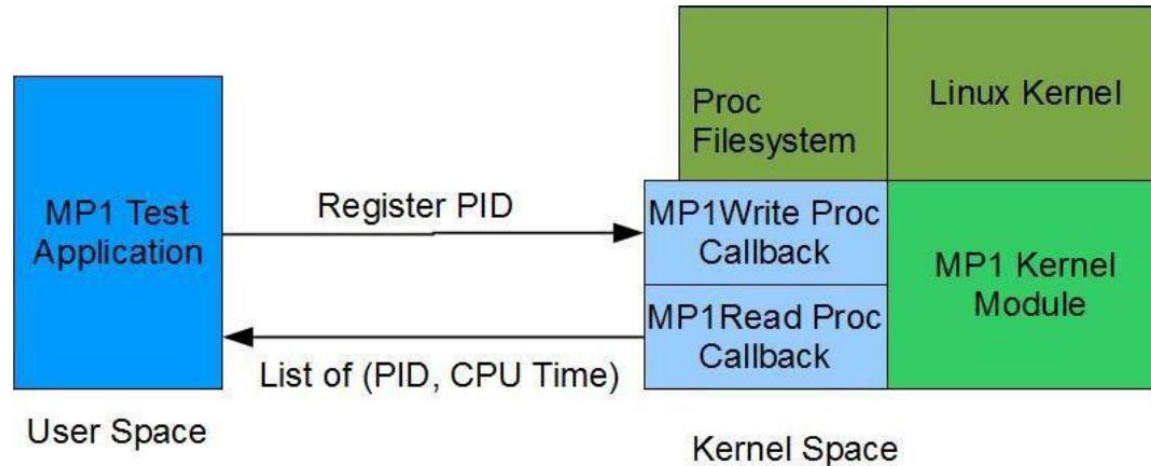
illinois.edu



Purpose of MP1

- Get yourself familiar with Linux kernel programming
- Learn to use kernel linked list data structure
- Learn to use /proc to communicate between kernel and user space program
- Timer, workqueue, interrupt etc.

MP1 Overview



- Build kernel module measure user app cpu time
- Use /proc file system to communicate between user program and kernel module
 - /proc/mp1/status
- Two-halves interrupt handler implementation
 - Top-half: interrupt handler
 - Bottom half: workqueue + worker thread

Kernel vs Application Programming

Kernel

- **No memory protection**
 - Share memory with devices, scheduler
 - Easily crash the system, very hard to debug
- **Sometimes no preemption**
 - Can hog the CPU
 - Concurrency is hard
- **No libraries**
 - No printf, fopen
- **No access to files**
- **Direct access to hardware**

Application

- **Memory protection**
 - Segmentation fault
 - Can conveniently debug the program
- **Preemption**
 - Scheduling is not our responsibility
- **Signals (Control-C)**
- **Libraries**
- **In Linux, everything is a file descriptor**
- **Access to hardware as files**

Linux Kernel Module (LKM)

- LKM are pieces of code that can be loaded and unloaded into the kernel upon demand
 - No need to modify the kernel source code for all MPs in this course
- Separate compilation
- Runtime linkage
- Entry and Exit functions

```
#include <linux/module.h>
#include <linux/kernel.h>

static int __init myinit(void)
{ printk(KERN_ALERT "Hello, world\n"); return 0;
}

static void __exit myexit(void)
{ printk(KERN_ALERT "Goodbye, World\n");
}

module_init(myinit);
module_exit(myexit);
MODULE_LICENSE("GPL");
```

LKM “Hello World”

```
#include <linux/module.h>
#include <linux/kernel.h>
static int __init myinit(void)
{
    printk(KERN_ALERT "Hello, world\n");
    return 0;
}
static void __exit myexit(void)
{
    printk(KERN_ALERT "Goodbye, World\n");
}
module_init(myinit);
module_exit(myexit);
MODULE_LICENSE("GPL");
```

- Edit source file as above

LKM “Hello World”

```
obj-m += hello.o
all:
    make -C /lib/modules/$(shell uname -r)/build M=$(PWD) modules
clean:
    make -C /lib/modules/$(shell uname -r)/build M=$(PWD) clean
```

- Edit the Makefile
- For MP1, the Makefile is provided
 - It can be reused for MP2/MP3

LKM “Hello World”

```
cs423@cs423-vm: ~/cs423/demo/mp1
File Edit View Search Terminal Help
cs423@cs423-vm:~/cs423/demo/mp1$ vim Makefile
cs423@cs423-vm:~/cs423/demo/mp1$ make
make -C /lib/modules/3.13.0-44-generic/build M=/home/cs423/cs423/demo/mp1 modules
make[1]: Entering directory `/usr/src/linux-headers-3.13.0-44-generic'
  CC [M]  /home/cs423/cs423/demo/mp1/hello.o
  Building modules, stage 2.
  MODPOST 1 modules
  CC      /home/cs423/cs423/demo/mp1/hello.mod.o
  LD [M]  /home/cs423/cs423/demo/mp1/hello.ko
make[1]: Leaving directory `/usr/src/linux-headers-3.13.0-44-generic'
cs423@cs423-vm:~/cs423/demo/mp1$
```

- Make
 - Compile the module

LKM “Hello World”

```
cs423@cs423-vm: ~/cs423/demo/mp1
File Edit View Search Terminal Help
cs423@cs423-vm:~/cs423/demo/mp1$ vim Makefile
cs423@cs423-vm:~/cs423/demo/mp1$ make
make -C /lib/modules/3.13.0-44-generic/build M=/home/cs423/cs423/demo/mp1 modules
make[1]: Entering directory `/usr/src/linux-headers-3.13.0-44-generic'
  CC [M]  /home/cs423/cs423/demo/mp1/hello.o
  Building modules, stage 2.
  MODPOST 1 modules
  CC      /home/cs423/cs423/demo/mp1/hello.mod.o
  LD [M]  /home/cs423/cs423/demo/mp1/hello.ko
make[1]: Leaving directory `/usr/src/linux-headers-3.13.0-44-generic'
cs423@cs423-vm:~/cs423/demo/mp1$ ls
hello.c  hello.ko  hello.mod.c  hello.mod.o  hello.o  Makefile  modules.order  Module.symvers
cs423@cs423-vm:~/cs423/demo/mp1$
```

- ls
 - Module is compiled as hello.ko

LKM “Hello World”

```
cs423@cs423-vm:~/cs423/demo/mp1$ ls
hello.c hello.ko hello.mod.c hello.mod.o hello.o Makefile modules.order Module.symvers
cs423@cs423-vm:~/cs423/demo/mp1$ sudo insmod hello.ko
[sudo] password for cs423:
cs423@cs423-vm:~/cs423/demo/mp1$ lsmod
Module                Size  Used by
hello                  12421  0
```

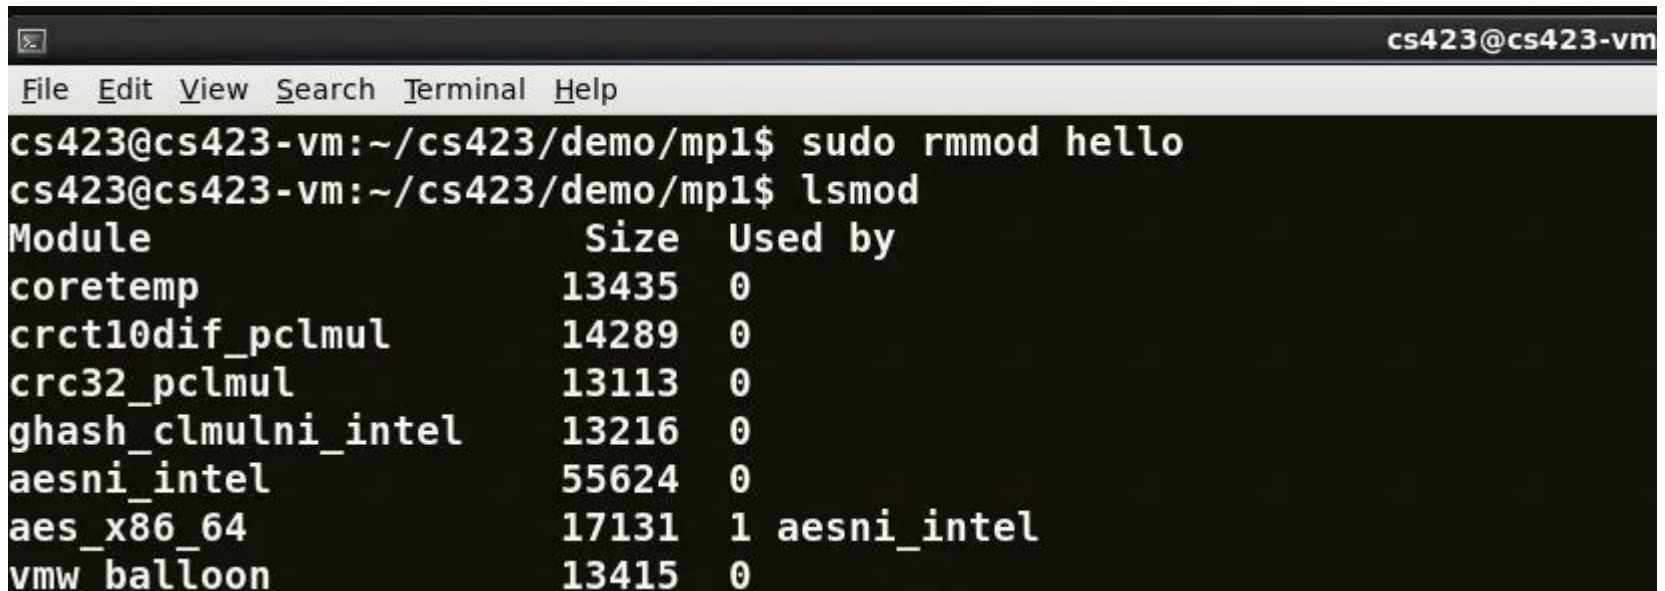
- `sudo insmod hello.ko`
 - install the module
- `lsmod`
 - You can see the hello module on the top of the list

LKM “Hello World”

```
cs423@cs423-vm:~/cs423/demo/mp1$ modinfo hello.ko
filename:          /home/cs423/cs423/demo/mp1/hello.ko
license:           GPL
srcversion:        0D371D51CDEEAE5E55A3841
depends:
vermagic:          3.13.0-44-generic SMP mod_unload modversions
cs423@cs423-vm:~/cs423/demo/mp1$
```

- modinfo
 - Check the module information

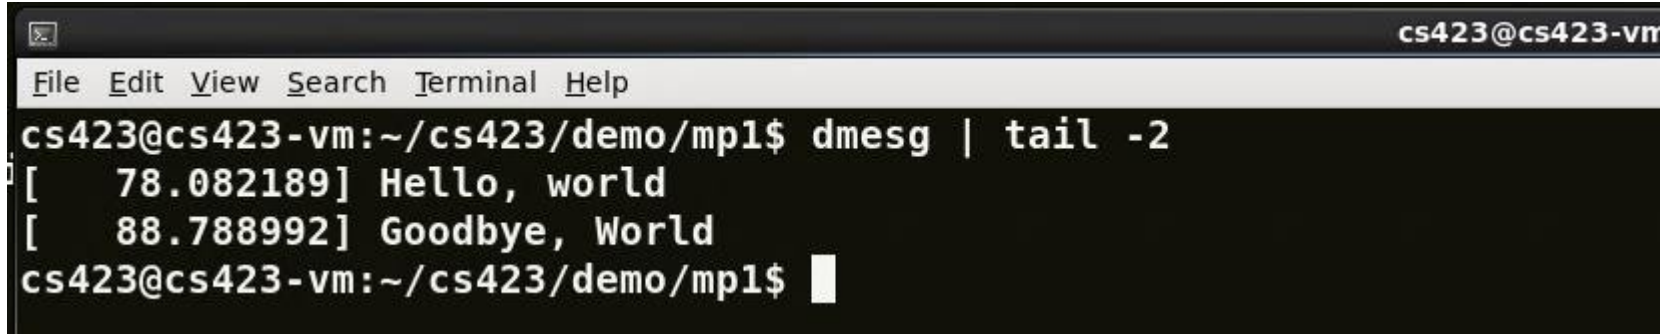
LKM “Hello World”



```
cs423@cs423-vm:~/cs423/demo/mp1$ sudo rmmod hello
cs423@cs423-vm:~/cs423/demo/mp1$ lsmod
Module                Size  Used by
coretemp              13435  0
crct10dif_pclmul      14289  0
crc32_pclmul          13113  0
ghash_clmulni_intel   13216  0
aesni_intel           55624  0
aes_x86_64            17131  1 aesni_intel
vmw_balloon           13415  0
```

- `sudo rmmod hello`
 - uninstall the module

LKM “Hello World”

A terminal window with a dark background and light text. The title bar at the top right says 'cs423@cs423-vm'. Below the title bar is a menu bar with 'File', 'Edit', 'View', 'Search', 'Terminal', and 'Help'. The terminal content shows a command prompt 'cs423@cs423-vm:~/cs423/demo/mp1\$' followed by the command 'dmesg | tail -2'. The output consists of two lines: '[78.082189] Hello, world' and '[88.788992] Goodbye, World'. The prompt is followed by a cursor.

```
cs423@cs423-vm:~/cs423/demo/mp1$ dmesg | tail -2
[ 78.082189] Hello, world
[ 88.788992] Goodbye, World
cs423@cs423-vm:~/cs423/demo/mp1$
```

- dmesg
 - Check kernel messages (printk)
 - Very useful to debug the module
 - dmesg | tail -n
 - Check the last n lines of kernel message

LKM “Hello World”

- To summarize
 - `sudo insmod hello.ko`
 - install the kernel module
 - `lsmod`
 - Check if the module is loaded
 - All loaded modules can be found `/proc/modules`
 - `sudo rmmod hello`
 - Unload the module

Kernel Module vs Application Programming

Kernel Module (LKM)

- Start with `module_init()`
 - Set up the kernel
- Runs in kernel space
- The module does nothing until one of the module functions are called by the kernel
- Ends with `module_exit()`

Applications

- Start with `main()`
- Runs in user space
- Executes a bunch of instructions
- Terminates

Functions available to LKM

- Applications have access to library functions
 - `printf()`, `malloc()`, `free()`
- Kernel modules do not have access to library functions except that provided by kernel
 - `printk()`, `kmalloc()`, `kfree()`, `vmalloc()`
 - Check `/proc/kallsyms` to see a list of kernel provided functions
- Check Linux Kernel Programming Guide page and references on the MP1 page

The /proc file system

- /proc is a virtual file system that allow communication between kernel and use space
- It doesn't contain 'real' files but runtime system information
 - system memory, devices mounted, hardware configuration
- Widely used for many reportings
 - /proc/modules, /proc/meminfo, /proc/cpuinfo

The /proc file system

```
cs423@cs423-vm:/proc$ ls
1      1329 1453 31    7221  execdomains  pagetypeinfo
10     133   1457 311   7319  fb           partitions
1052   134   146  315   7384  filesystems  sched_debug
1089   135   147  32    764   fs          schedstat
11     1351  148  33    8     interrupts  scsi
115    1354  1488 401   830   iomem       self
1153   136   15   444   832   ioports     slabinfo
1154   137   16   445   9     ipmi        softirqs
116    138   17   45   918   irq         stat
117    1383  174   460   923   kallsyms    swaps
118    139   177   47   928   kcore       sys
119    1397  18   48   929   keys        sysrq-trigger
12     1398  189   480   932   key-users   sysvipc
120    14   19   5     970   kmsg        timer_list
121    140   190   500   986   kpagecount  timer_stats
122    1404  2     502   ACPI       tty
123    1405  20    524   buddyinfo  uptime
124    141   21    644   bus        version
125    1410  22    69   cgroups   version_signature
126    1414  23    7     cmdline   vmallocinfo
127    1417  24    70    consoles  vmstat
128    142   25   7081  cpuinfo    zoneinfo
129    1421  27   7082  crypto     modules
13     1425  28   7084  devices    mounts
130    143   29   7184  diskstats  mpt
131    144   3     722   dma        mtrr
132    145   30   7220  driver     net
cs423@cs423-vm:/proc$
```

The /proc file system

```
cs423@cs423-vm:/proc$ cat /proc/cpuinfo
processor       : 0
vendor_id      : GenuineIntel
cpu family     : 6
model          : 62
model name     : Intel(R) Xeon(R) CPU E5-2670 v2 @ 2.50GHz
stepping       : 4
microcode      : 0x427
cpu MHz        : 2500.000
cache size     : 25600 KB
physical id    : 0
siblings       : 1
core id        : 0
cpu cores      : 1
apicid         : 0
initial apicid : 0
fpu            : yes
fpu_exception  : yes
cpuid level    : 13
wp             : yes
flags          : fpu vme de pse tsc msr pae mce cx8 apic sep mtrr pge mca cmov pat pse36 clflush
                dts mmx fxsr sse sse2 ss syscall nx rdtscp lm constant_tsc arch_perfmon pebs bts nopl xtopology
tsc_reliable   nonstop_tsc aperfmperf eagerfpu pni pclmulqdq ssse3 cx16 pcid sse4_1 sse4_2 x2apic p
opcnt aes xsave avx f16c rdrand hypervisor lahf_lm ida arat xsaveopt pln pts dtherm fsgsbase smep
bogomips       : 5000.00
clflush size   : 64
cache_alignment : 64
address sizes   : 40 bits physical, 48 bits virtual
power management:
```

The /proc file system

```
cs423@cs423-vm:/proc$ cat /proc/meminfo
MemTotal:      1017836 kB
MemFree:       422048 kB
Buffers:       68584 kB
Cached:        383060 kB
SwapCached:    0 kB
Active:        236344 kB
Inactive:      276500 kB
Active(anon):   61836 kB
Inactive(anon): 4088 kB
Active(file):   174508 kB
Inactive(file): 272412 kB
Unevictable:    0 kB
Mlocked:        0 kB
SwapTotal:     1046524 kB
SwapFree:      1046524 kB
Dirty:         24 kB
Writeback:      0 kB
AnonPages:     61196 kB
Mapped:        33832 kB
Shmem:         4728 kB
Slab:          46440 kB
SReclaimable:  33392 kB
SUnreclaim:    13048 kB
KernelStack:   1712 kB
PageTables:    5976 kB
NFS_Unstable:  0 kB
Bounce:        0 kB
```

Using /proc in MP1

```
19 extern struct proc_dir_entry *proc_mkdir(const char *, struct proc_dir_entry *);
```

- Create a directory under /proc
 - `proc_mkdir()`

```
30 static inline struct proc_dir_entry *proc_create(  
31     const char *name, umode_t mode, struct proc_dir_entry *parent,  
32     const struct file_operations *proc_fops)  
33 {  
34     return proc_create_data(name, mode, parent, proc_fops, NULL);  
35 }  
36
```

- Create a file under /proc
 - `proc_create()`

Using /proc in MP1

```
1486 struct file_operations {
1487     struct module *owner;
1488     loff_t (*llseek) (struct file *, loff_t, int);
1489     ssize_t (*read) (struct file *, char __user *, size_t, loff_t *);
1490     ssize_t (*write) (struct file *, const char __user *, size_t, loff_t *);
1491     ssize_t (*aio_read) (struct kiocb *, const struct iovec *, unsigned long, loff_t);
1492     ssize_t (*aio_write) (struct kiocb *, const struct iovec *, unsigned long, loff_t);
1493     ssize_t (*read_iter) (struct kiocb *, struct iov_iter *);
1494     ssize_t (*write_iter) (struct kiocb *, struct iov_iter *);
1495     int (*iterate) (struct file *, struct dir_context *);
1496     unsigned int (*poll) (struct file *, struct poll_table_struct *);
1497     long (*unlocked_ioctl) (struct file *, unsigned int, unsigned long);
1498     long (*compat_ioctl) (struct file *, unsigned int, unsigned long);
1499     int (*mmap) (struct file *, struct vm_area_struct *);
1500     int (*open) (struct inode *, struct file *);
1501     int (*flush) (struct file *, fl_owner_t id);
1502     int (*release) (struct inode *, struct file *);
1503     int (*fsync) (struct file *, loff_t, loff_t, int datasync);
1504     int (*aio_fsync) (struct kiocb *, int datasync);
1505     int (*fasync) (int, struct file *, int);
1506     int (*lock) (struct file *, int, struct file_lock *);
1507     ssize_t (*sendpage) (struct file *, struct page *, int, size_t, loff_t *, int);
1508     unsigned long (*get_unmapped_area) (struct file *, unsigned long, unsigned long, unsigned long, unsig
1509     int (*check_flags) (int);
1510     int (*flock) (struct file *, int, struct file_lock *);
1511     ssize_t (*splice_write) (struct pipe_inode_info *, struct file *, loff_t *, size_t, unsigned int);
1512     ssize_t (*splice_read) (struct file *, loff_t *, struct pipe_inode_info *, size_t, unsigned int);
1513     int (*setlease) (struct file *, long, struct file_lock **, void **);
1514     long (*fallocate) (struct file *file, int mode, loff_t offset,
1515                        loff_t len);
1516     int (*show_fdinfo) (struct seq_file *m, struct file *f);
1517 };
```


Using /proc in MP1

Sample code:

```
#define FILENAME "status"
#define DIRECTORY "mp1"
static struct proc_dir_entry *proc_dir;
static struct proc_dir_entry *proc_entry;
static ssize_t mp1_read (struct file *file, char __user *buffer, size_t count, loff_t *data){
    // implementation goes here...
}
static ssize_t mp1_write (struct file *file, const char __user *buffer, size_t count, loff_t *data){
    // implementation goes here...
}
static const struct file_operations mp1_file = {
    .owner = THIS_MODULE,
    .read  = mp1_read,
    .write = mp1_write,
};
int __init mp1_init(void){
    proc_dir = proc_mkdir(DIRECTORY, NULL);
    proc_entry = proc_create(FILENAME, 0666, proc_dir, & mp1_file);
}
```

Using /proc in MP1

- Within mp1_read/mp1_write, you may need to move data between kernel/user space
 - Copy_from_user()
 - Copy_to_user()

Sample code (There are other ways of implementing it):

```
static ssize_t mp1_read (struct file *file, char __user *buffer, size_t count, loff_t *data){
    // implementation goes here...
    int copied;
    char * buf;
    buf = (char *) kmalloc(count,GFP_KERNEL);
    copied = 0;
    //... put something into the buf, updated copied
    copy_to_user(buffer, buf, copied);
    kfree(buf);
    return copied ;
}
```

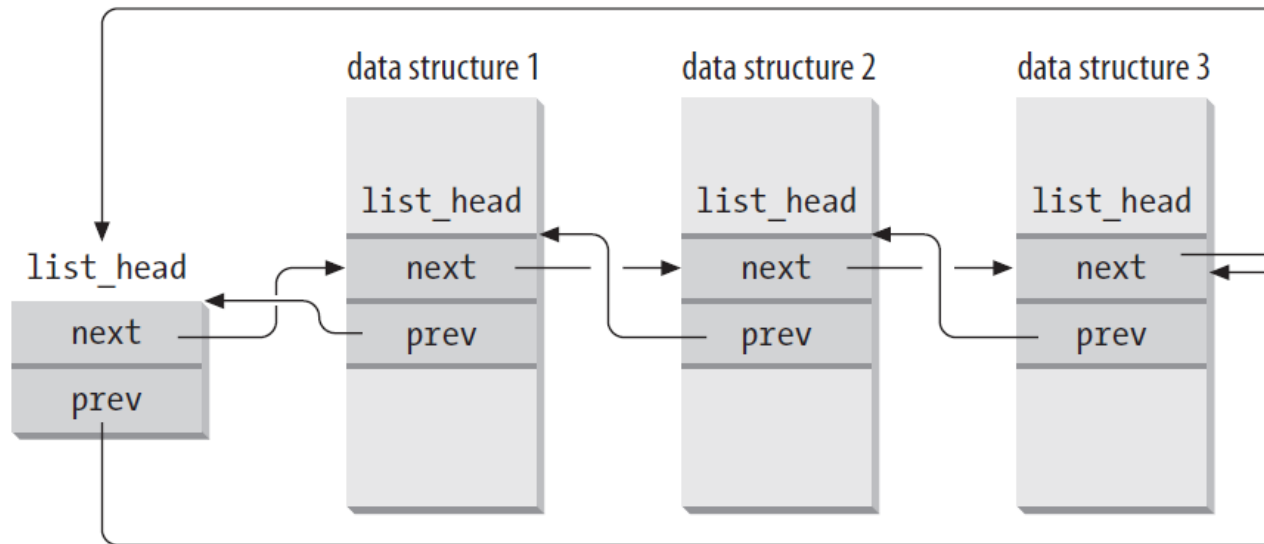
Linux Kernel List

- You will use Linux list to store all registered user processes
- Linux kernel list is a widely used data structure in Linux kernel
 - Define in <linux/linux.h>
 - You MUST get familiar of how to use it
 - Can be used as follows


```
struct list_head{  
    struct list_head *next;  
    struct list_head *prev;  
};
```

```
struct my_cool_list{  
    struct list_head list; /* kernel's list structure */  
    int my_cool_data;  
    void* my_cool_void;  
};
```

Linux Kernel List



(a) a doubly linked list with three elements



(b) an empty doubly linked list

Figure 3-3. Doubly linked lists built with `list_head` data structures

Linux Kernel List

- Some useful APIs
 - `LIST_HEAD(new_list)`
 - `list_add(struct list_head *new, struct list_head *head)`
 - `list_for_each_safe(pos, n, head)`
 - `list_entry(ptr, type, member)`
 - `list_del(pos)`
 - `list_for_each_entry(pos, head, member)`
 - `List_empty(ptr)`

Kernel Timer

- Operate in jiffies domain
 - msec_to_jiffies() to convert ms to jiffies
 - jiffies_to_msec() to convert jiffies to ms

```
struct timer_list {  
    /* ... */  
    unsigned long expires;  
    void (*function)(unsigned long);  
    unsigned long data;  
};
```

- The expires field represents the jiffies value when the timer is expected to run

Kernel Timer

- Some useful API
 - `void setup_timer(struct timer_list *timer, void(*function)(unsigned long), unsigned long data)`
 - `int mod_timer(struct timer_list *timer, unsigned long expires)`
 - `void del_timer(struct timer_list *timer)`
 - `void init_timer(struct timer_list *timer);`
 - `struct timer_list TIMER_INITIALIZER(_function, _expires, _data);`
 - `void add_timer(struct timer_list * timer);`

Workqueue

- Allow kernel code to request that a function be called at some future time
 - Workqueue functions can sleep
 - Can be used to implement to bottom half of the interrupt handlers
- Some useful API
 - `INIT_WORK (struct work_struct *work, void (*function)(void *), void *data)`
 - `void flush_workqueue (struct workqueue_struct *queue)`
 - `void destroy_workqueue (struct workqueue_struct *queue)`
 - `int queue_work (struct workqueue_struct *queue, struct work_struct *work)`

More questions?

- Office hours
- Piazza